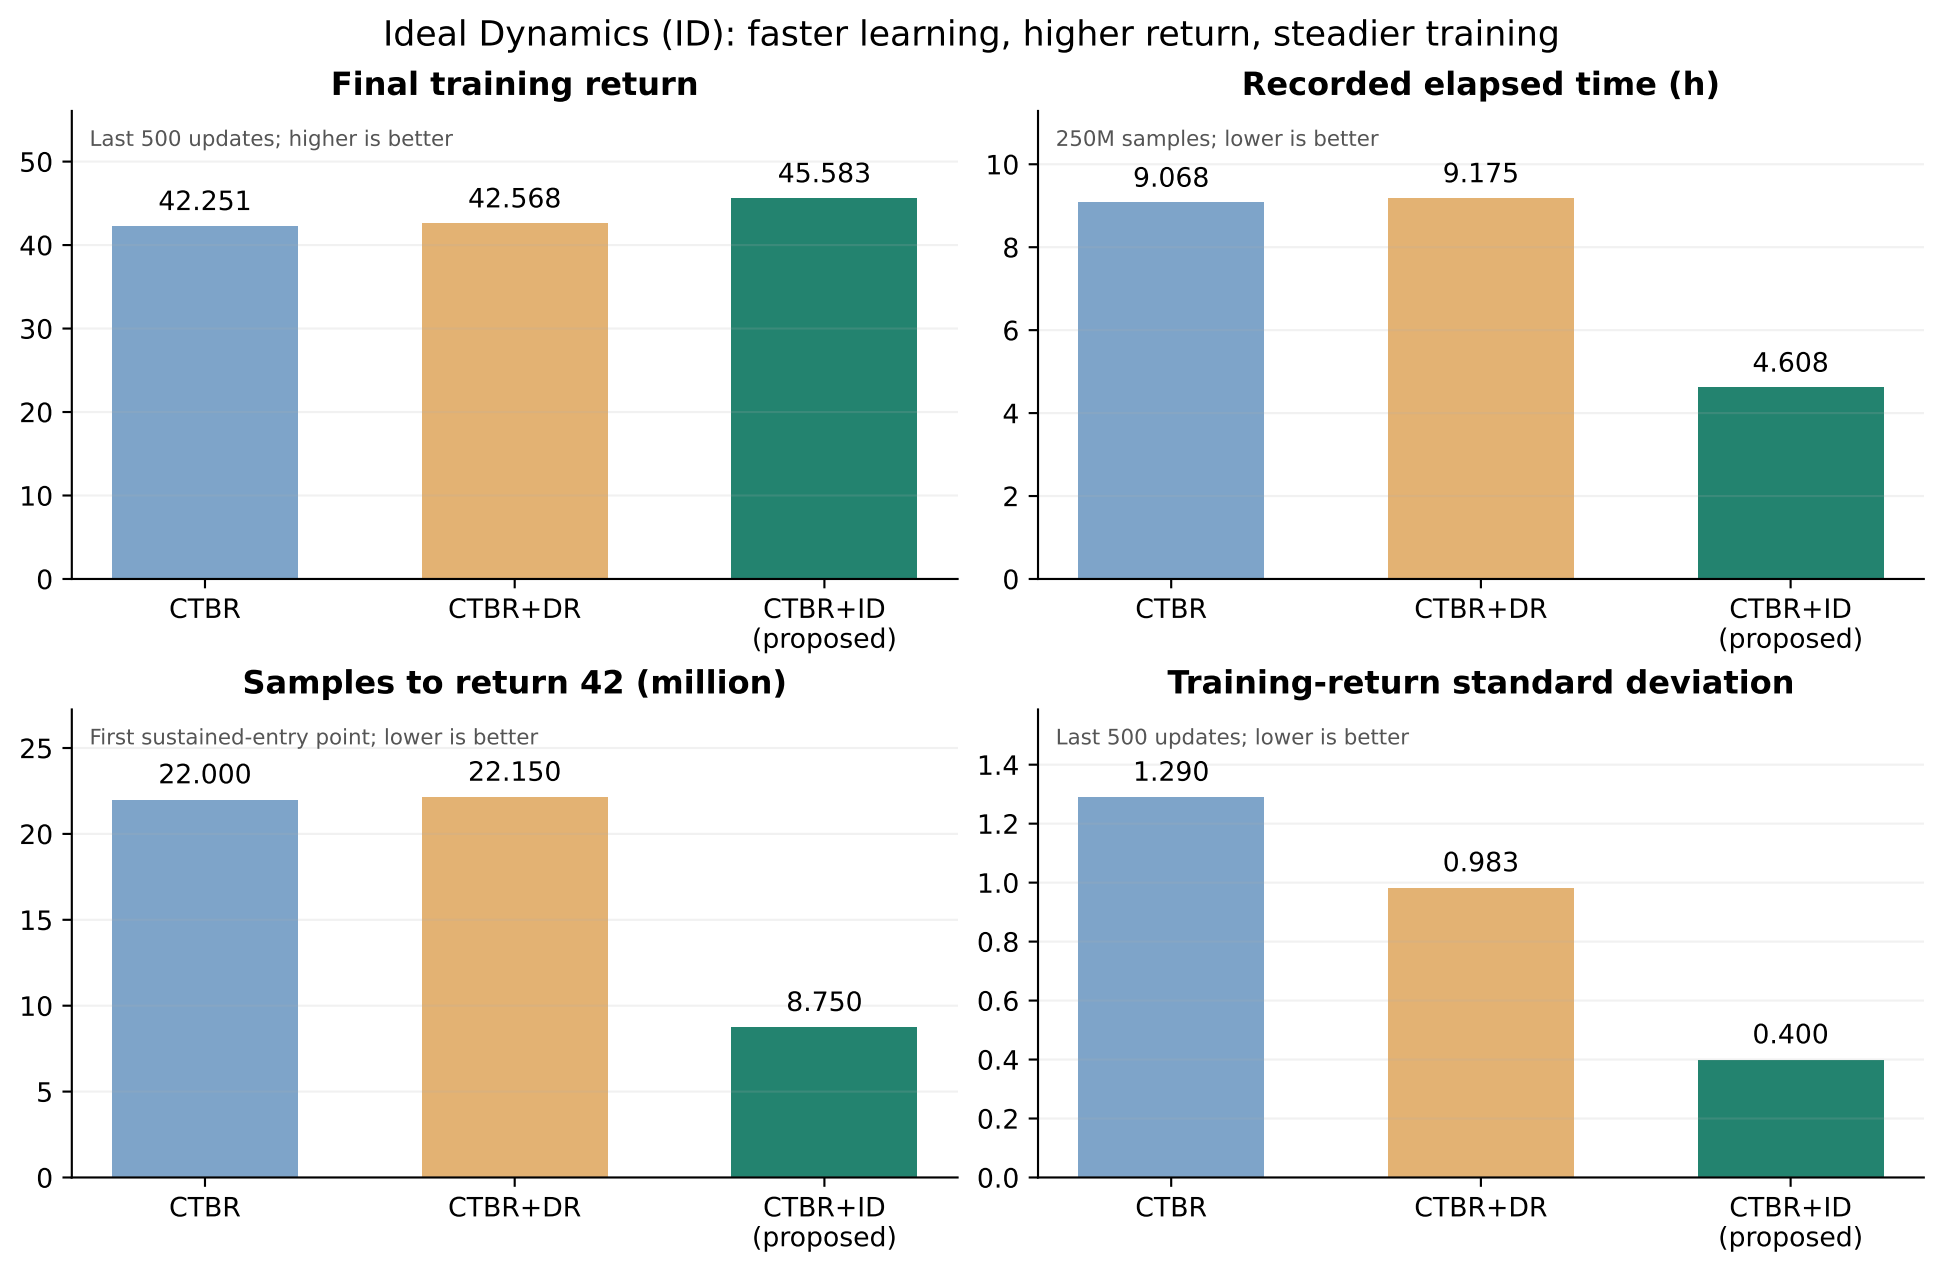
Ideal Dynamics (ID): faster learning, higher return, steadier training
Final training return
Last 500 updates; higher is better
42.251
42.568
45.583
CTBR
CTBR+DR
CTBR+ID (proposed)
Recorded elapsed time (h)
250M samples; lower is better
9.068
9.175
4.608
Samples to return 42 (million)
First sustained-entry point; lower is better
22.000
22.150
8.750
Training-return standard deviation
Last 500 updates; lower is better
1.290
0.983
0.400