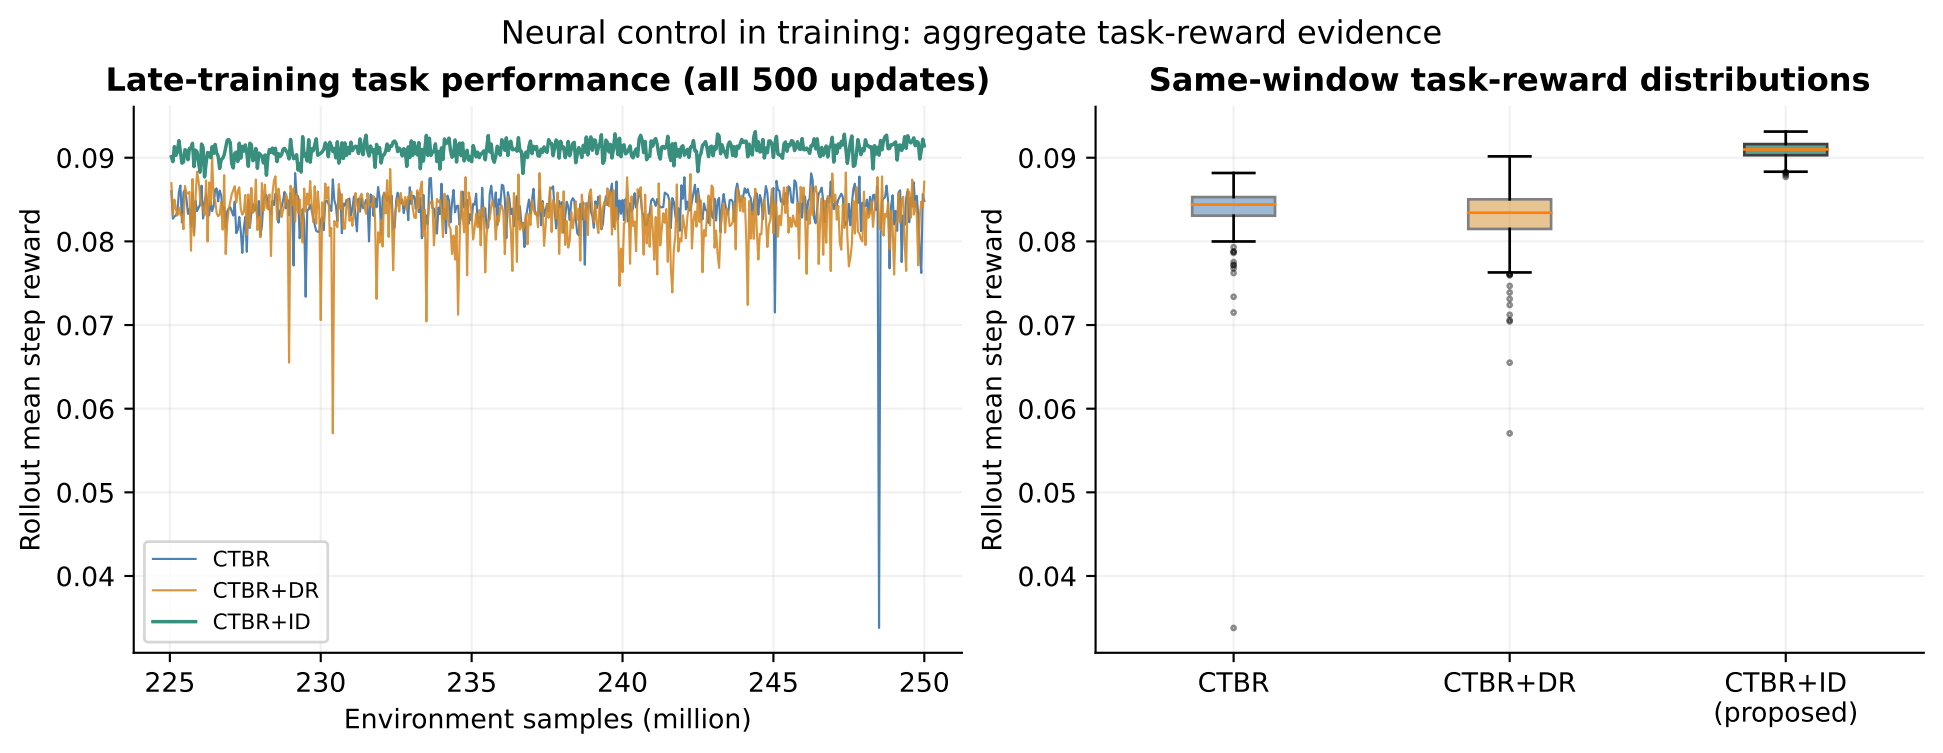
Neural control in training: aggregate task-reward evidence
Late-training task performance (all 500 updates)
Same-window task-reward distributions
Rollout mean step reward
0.09
0.08
0.07
0.06
0.05
0.04
CTBR
CTBR+DR
CTBR+ID
225
230
235
240
245
250
Environment samples (million)
CTBR
CTBR+DR
CTBR+ID
(proposed)